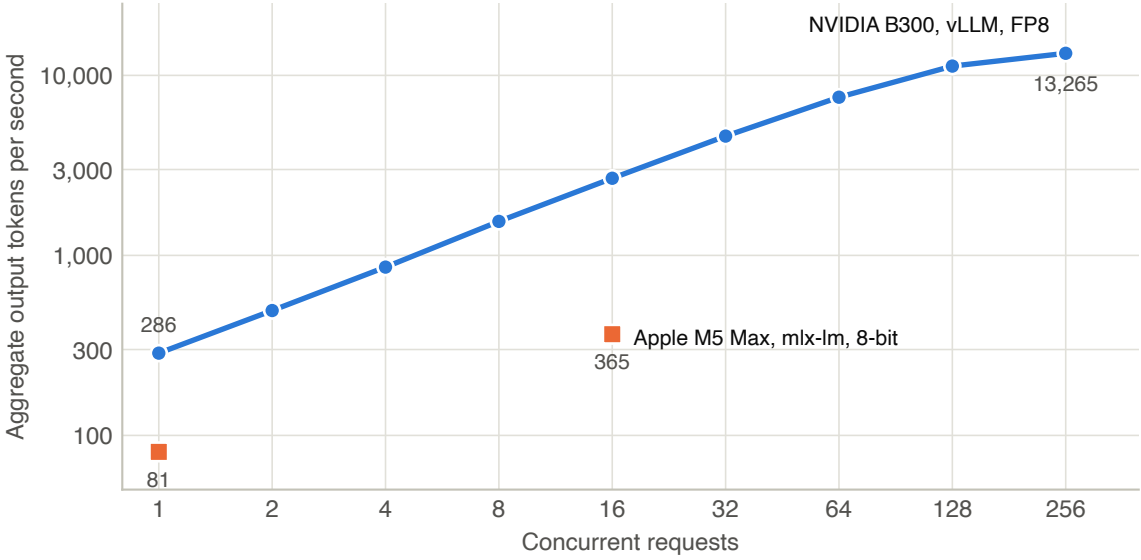
NVIDIA B300, vLLM, FP8
13,265
286
Apple M5 Max, mlx-lm, 8-bit
365
81
Aggregate output tokens per second
10,000
3,000
1,000
300
100
1
2
4
8
16
32
64
128
256
Concurrent requests